

Deploying at scale

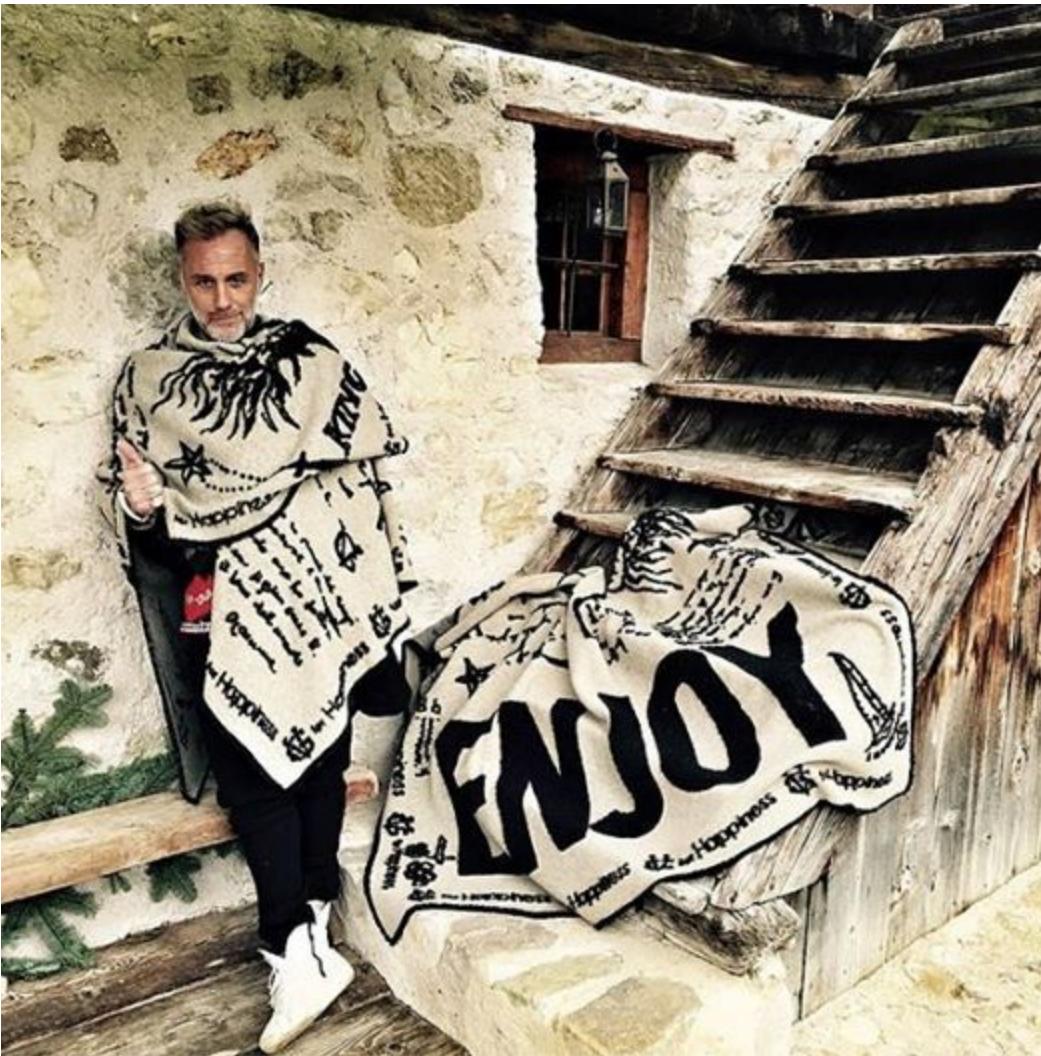
A developer journey

@lucamaraschi

With the “soft hand” touch of @matteocollina

DISCLAIMER

Nobody in this story, and no outfit or corporation, is based upon an actual person or outfit in the real world. The tale we are going to tell is completely fictional, but based upon the combined experiences of the speakers in building Node.js applications.



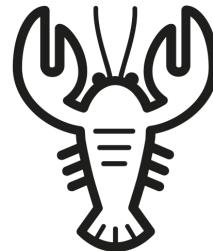


Created by Julia Holmberg
from Noun Project

Meet bob

He is a **.LAVA** developer

He writes applications on the **L2EE** stack, using **SQLobster** as a database.



Created by Rutmer Zijlstra
from Noun Project



Created by Julia Holmberg
from Noun Project

Bob reads:



ABOUT US SERVICES OPEN SOURCE EVENTS NODECRUNCH CAREERS CONTACT

NodeCrunch

Welcome to NodeCrunch, the nearForm blog. This is where we share information of all kinds, from Node.js news to new tool releases to our latest thinking on developing company values. We love to chat, so do post your comments below.



WHY NODE.JS IS BECOMING THE GO-TO TECHNOLOGY IN THE ENTERPRISE

By: Cian Ó Maidín

March 10, 2014 Feature articles, Node.js community Tagged Mail Online, node.js growth, PayPal, Voxer, WallMart, Yammer 65 Comments

RECENT POSTS

Node v6.0.0 released today
(April 26th)

Announcing microservices
day London 2016

Announcing NodeConf
London 2016

A Shiba Inu dog is sitting on a light-colored couch. It has a surprised or excited expression, with wide eyes and its mouth slightly open. The background shows a window with pink flowers and a small red sign.

Such speed

So enterprise

Wow JS



Created by Julia Holmberg
from Noun Project



Bob decides:

Let's rewrite my HUGE application in node.js
Using the MEAN stack



mongoDB®

mongoose

Express

node.js®

A large, light-colored rock formation with distinct horizontal layers, stands prominently against a vibrant sunset sky. The sky transitions from deep blue at the top to a warm orange and yellow near the horizon. In the foreground, the calm sea is visible with some rocky outcrops. A large, semi-transparent green ".js" logo is positioned in the lower right area of the rock formation.

.js

How do I deploy my Javascript monolith?



Created by Julia Holmberg
from Noun Project

I read StackOverflow!

SSH into the virtual machine

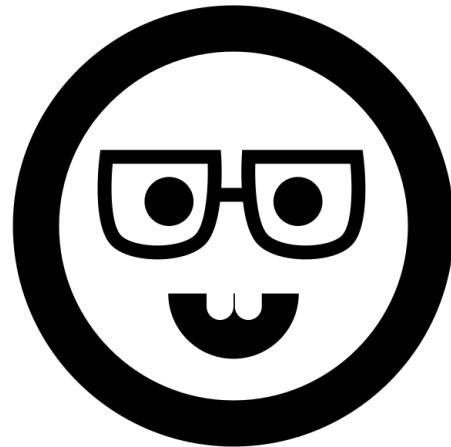
- snowflake servers
- git clone
- Forever or PM2!

No Integration tests, but 100%
code coverage

The new product has huge success

Bob is happy





And the performance fall

Bob is SAAD

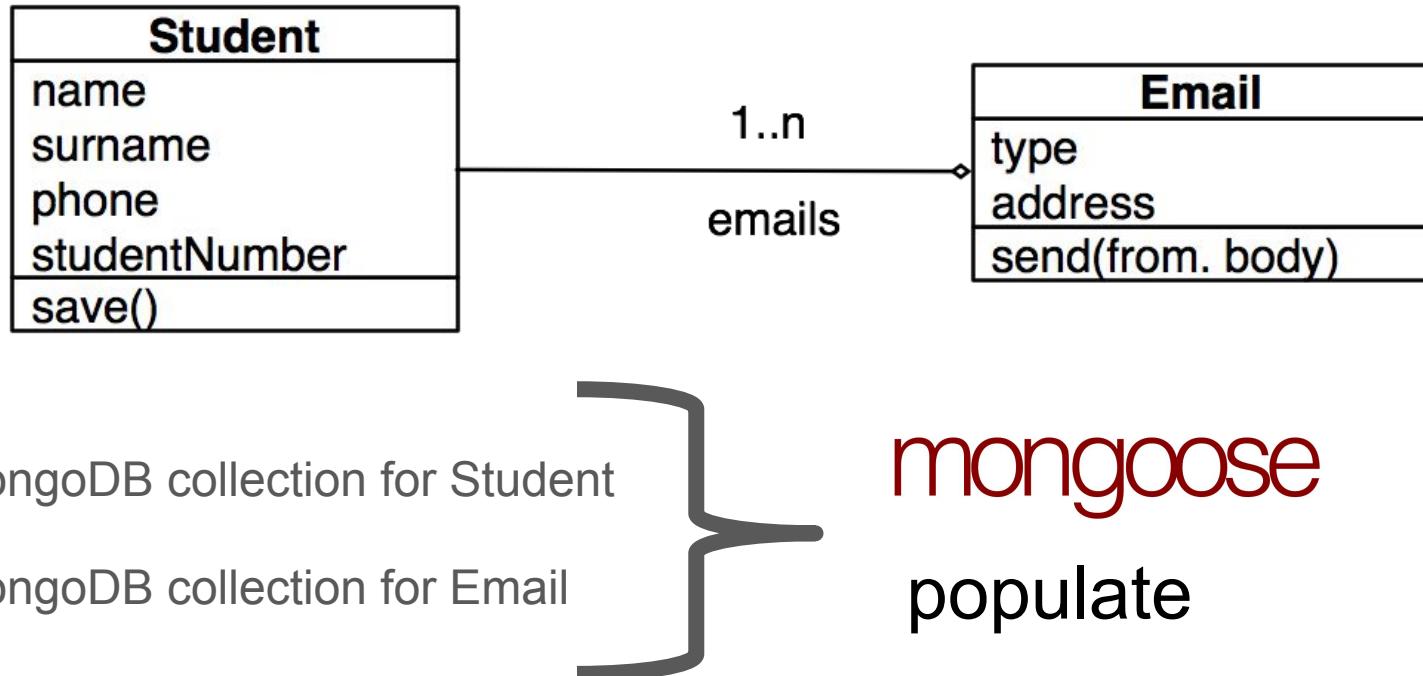
A Shiba Inu dog is sitting on a light-colored couch. It has a surprised or excited expression, with wide eyes and its mouth slightly open. The background shows a window with pink flowers and a small red sign.

Such speed

So enterprise

Wow JS

But my schema is relational...



Population

There are no joins in MongoDB but sometimes we still want references to documents in other collections. This is where population comes in.

Population is the process of automatically replacing the specified paths in the document with document(s) from other collection(s). We may populate a single document, multiple documents, plain object, multiple plain objects, or all objects returned from a query. Let's look at some examples.

```
var mongoose = require('mongoose')
, Schema = mongoose.Schema

var personSchema = Schema({
  _id      : Number,
  name    : String,
  age     : Number,
  stories : [{ type: Schema.Types.ObjectId, ref: 'Story' }]
});

var storySchema = Schema({
  _creator : { type: Number, ref: 'Person' },
  title    : String,
  fans     : [{ type: Number, ref: 'Person' }]
});

var Story  = mongoose.model('Story', storySchema);
var Person = mongoose.model('Person', personSchema);
```

Facts

MongoDB has no concept of relationship

We need relationships in our application

Let's build a JOIN engine inside our application!

Taken from the Mongoose website

Population

So far we haven't done anything much different. We've merely created a Person and a Story. Now let's take a look at populating our story's _creator using the query builder:

```
Story
.findOne({ title: 'Once upon a timex.' })
.populate('_creator')
.exec(function (err, story) {
  if (err) return handleError(err);
  console.log('The creator is %s', story._creator.name);
  // prints "The creator is Aaron"
});
```

Let's do TWO queries
to Mongo

If we used find() instead
of findOne(), we would
have done 1+N queries

Facts

MongoDB has no concept of relationship

We need relationships in our application

Let's build a JOIN
engine inside our
application!

NoSQL is hard!

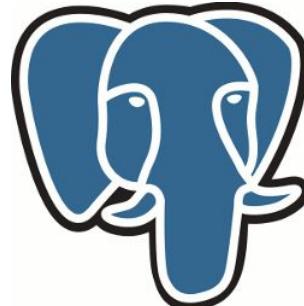
```
{  
  "name": "Matteo",  
  "surname": "Collina",  
  "phone": "+39 347123456",  
  "Description": "call me 24/7!!!"  
  "emails": [{  
    "type": "university",  
    "address":  
      "matteo.collina@cepu.org"  
  }]  
}
```

Relational vs. NoSQL

It is not Node fault!

Deduplication is accepted

Don't shoot in your own foot



PostgreSQL
the world's most advanced open source database

The fixed product has huge success

Bob is happy
again



...but my servers stop responding!!!!



A Shiba Inu dog is sitting on a light-colored couch. It has a surprised or excited expression, with wide eyes and its mouth slightly open. The background shows a window with pink flowers and a small red sign.

Such speed

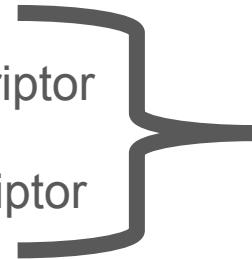
So enterprise

Wow JS

Exhausted file descriptors

Every incoming HTTP request is a file descriptor

Every outgoing HTTP request is a file descriptor



Fixed amount

1. You receive an HTTP request call to /faulty
2. In faulty, you call service A on /doSomething
3. You **pipe** the result of /doSomething to the HTTP response
4. An error happens, and **you do not close the HTTP response**

Exhausted file descriptors: solution

Replace
With

a.pipe(b)

pump(a, b)

★ pump public

pipe streams together and close all of them if one of them closes

pump is a small node module that pipes streams together and destroys all of them if one of them closes.

npm install pump

build passing

What problem does it solve?

When using standard source.pipe(dest) source will *not* be destroyed if dest emits close or an error. You are also not able to provide a callback to tell when then pipe has finished.

pump does these two things for you

We are back online!

Bob is happy
again
again



Created by Julia Holmberg
from Noun Project

But the Business (always) wants more...

- Scaling the development team from 1 (Bob) to 5
- Delivering more features in less time
- 1000 unit tests!
- Serve more customers!

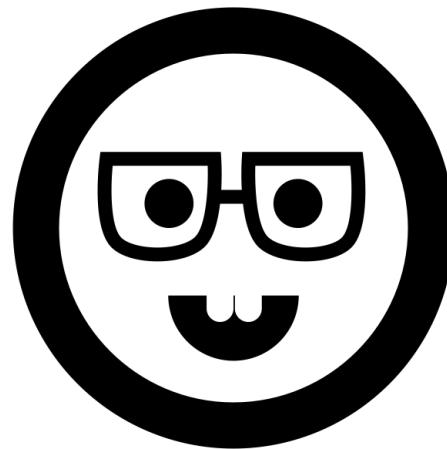


A Shiba Inu dog is sitting on a light-colored couch. It has a surprised or excited expression, with wide eyes and its mouth slightly open. The background shows a window with pink flowers and a small red sign.

Such speed

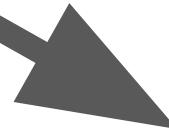
So enterprise

Wow JS



The new team does not deliver

Bob is SAAD



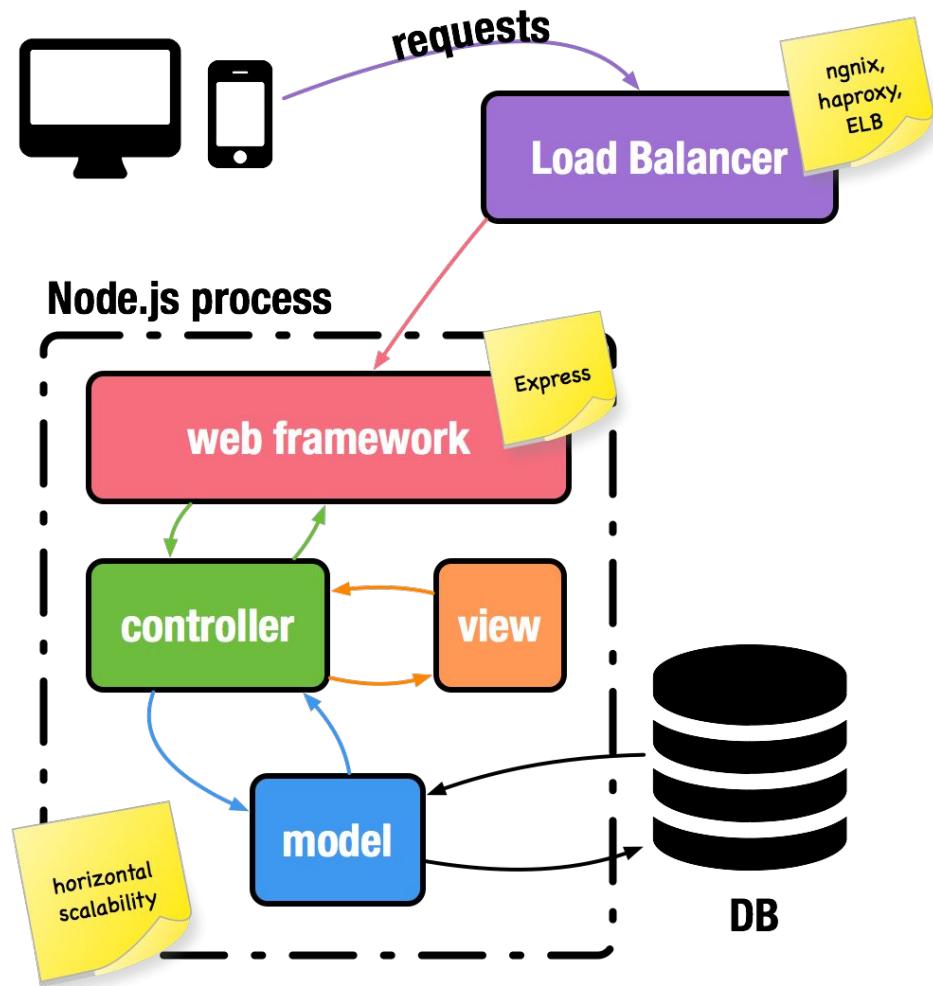
MVC architecture

Really monolithic

Adding a feature means changing all “layers” in an application

Does not scale with the team size

Equivalent VMs need to be dimensioned to handle the most demanding feature





Don't push anything
till tomorrow..

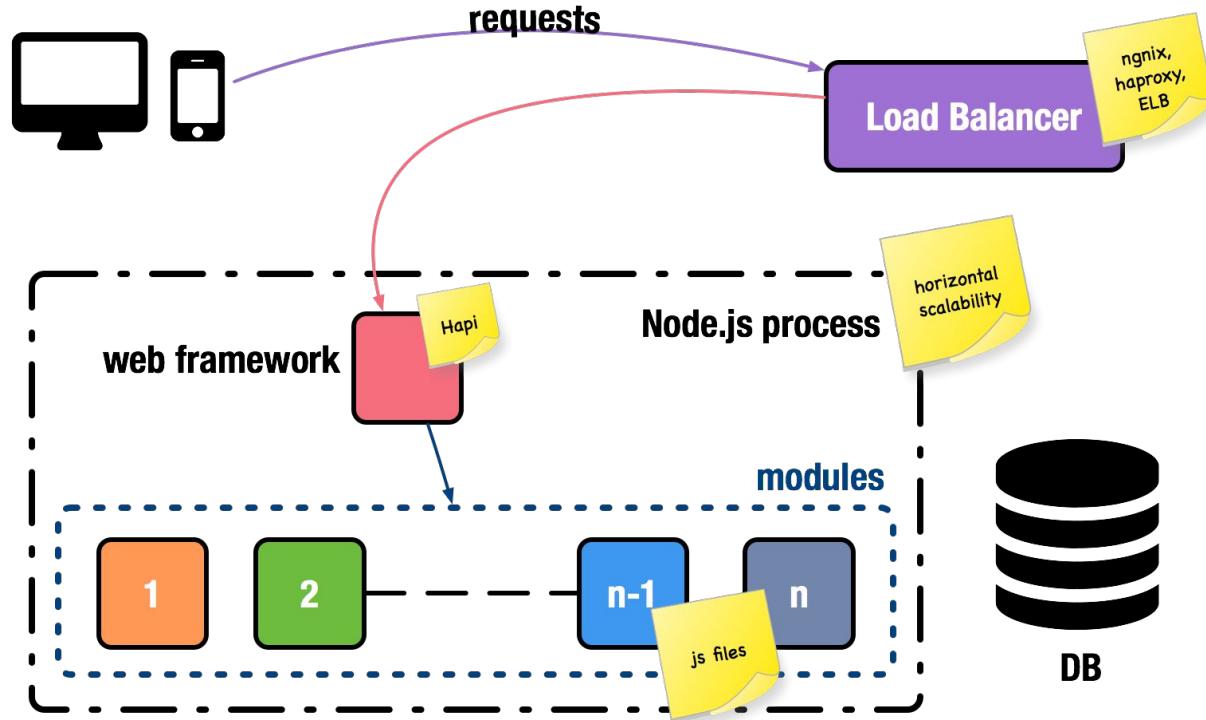
Developing with modules

Really monolithic

~~Adding a feature means changing all “layers” in an application~~

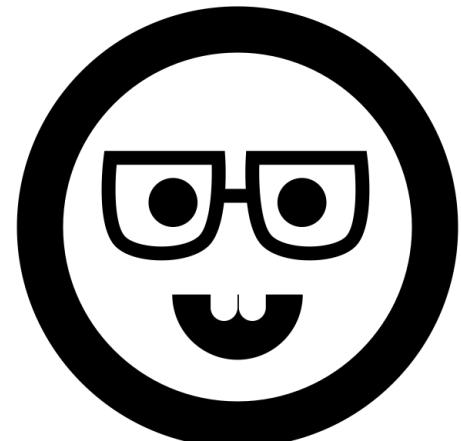
Does not scale with the team size *up to a certain point*

Equivalent VMs need to be dimensioned to handle the most demanding feature



Productivity goes up again!

Bob is happy
again
again
again

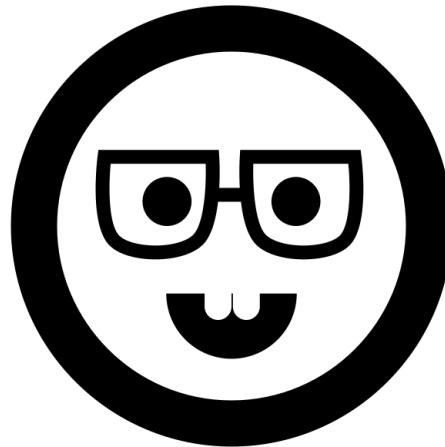


Created by Julia Holmberg
from Noun Project

But the Business (always) wants more...

- Scaling the development team from 5 to 20
- Delivering **even** more features in less time
- 10000 unit tests!
- Serve more customers!

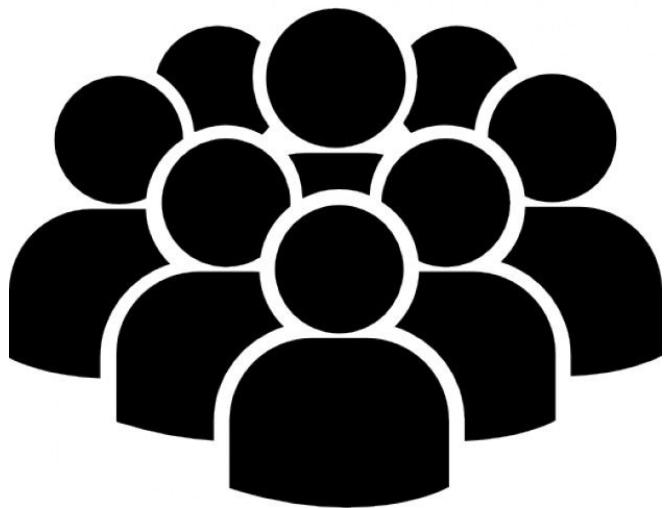




Bob is
DESPERATE

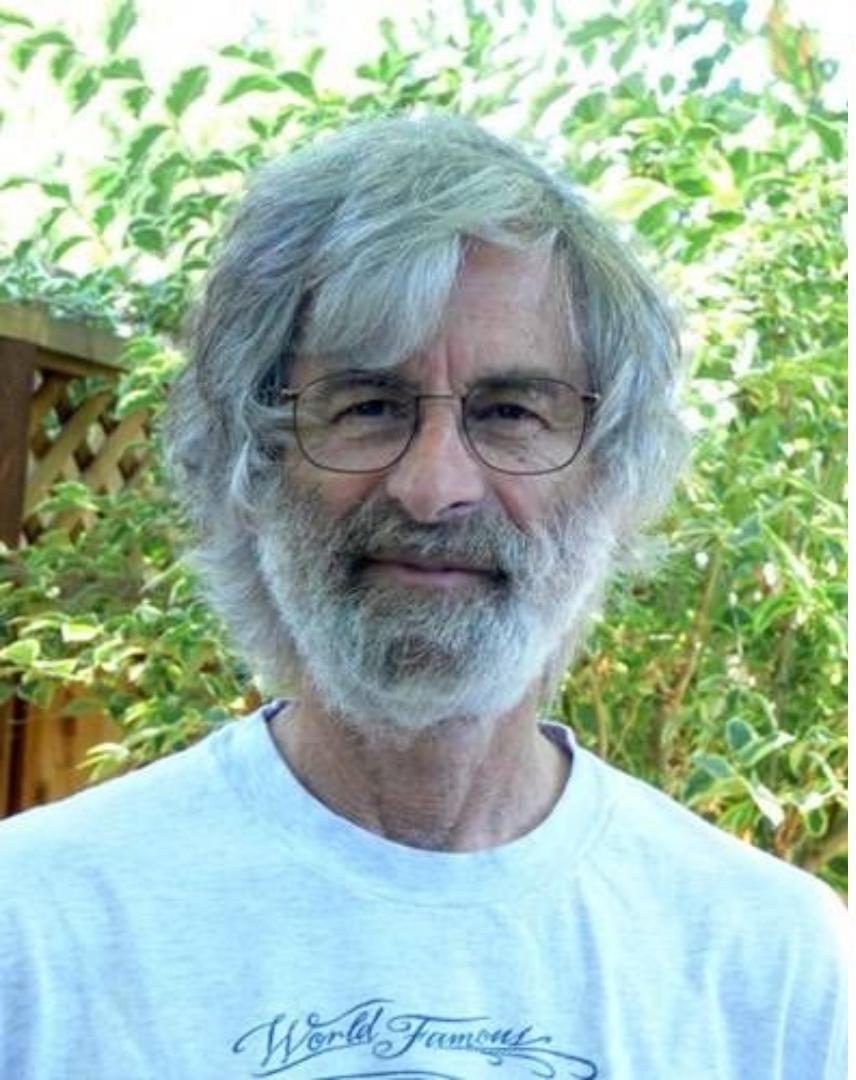
The NEW new team does not deliver



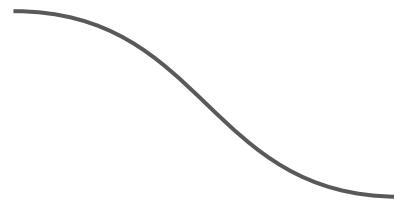




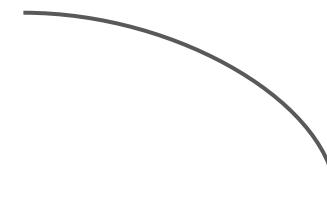
++
distribution



Missing Resiliency



&&



Fault tolerance

A pre-made “stack” might not be good



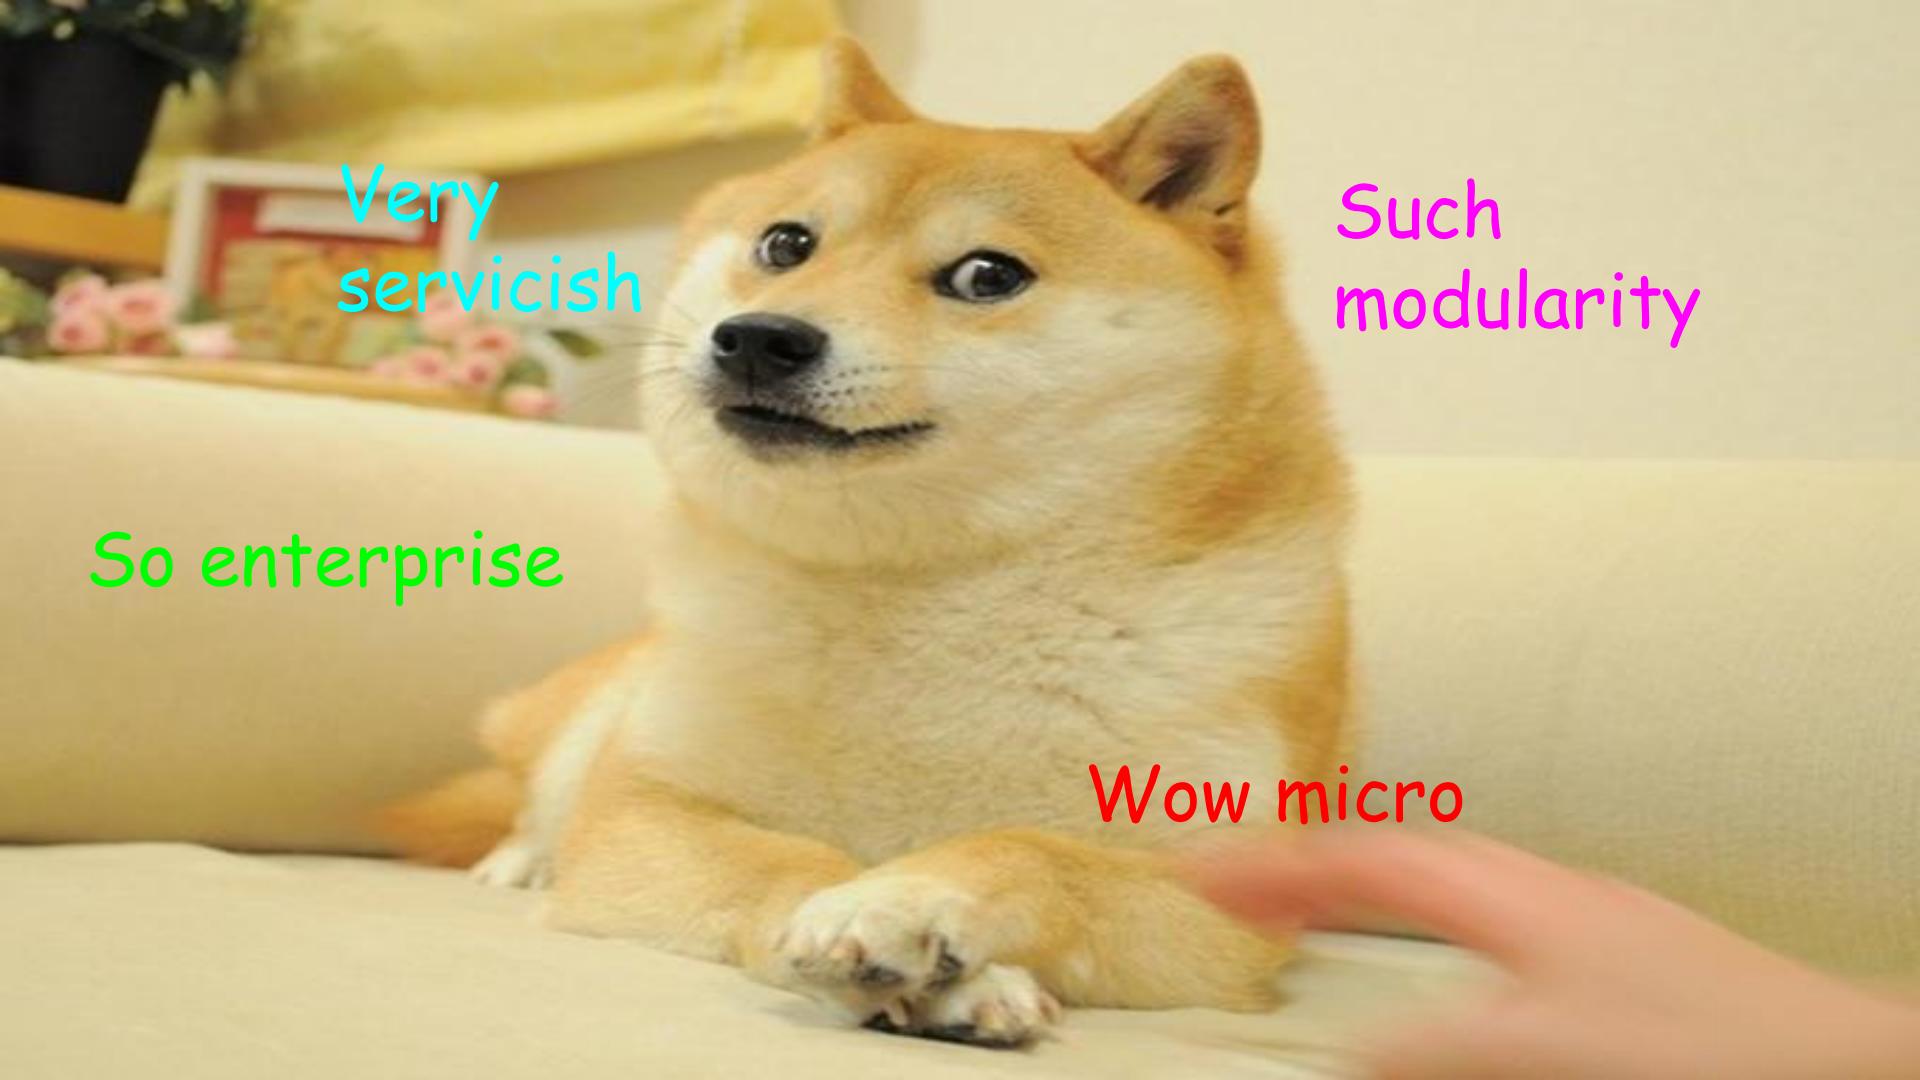


Let's throw
everything away...

But .LAVA42 does it better! ;-)



Hipsters use
microservices

A Shiba Inu dog is sitting on a light-colored couch. It has a white and tan coat. Its head is turned slightly to the left, looking towards the camera. The background shows a window with pink flowers outside.

So enterprise

Very
servicish

Wow micro

Such
modularity

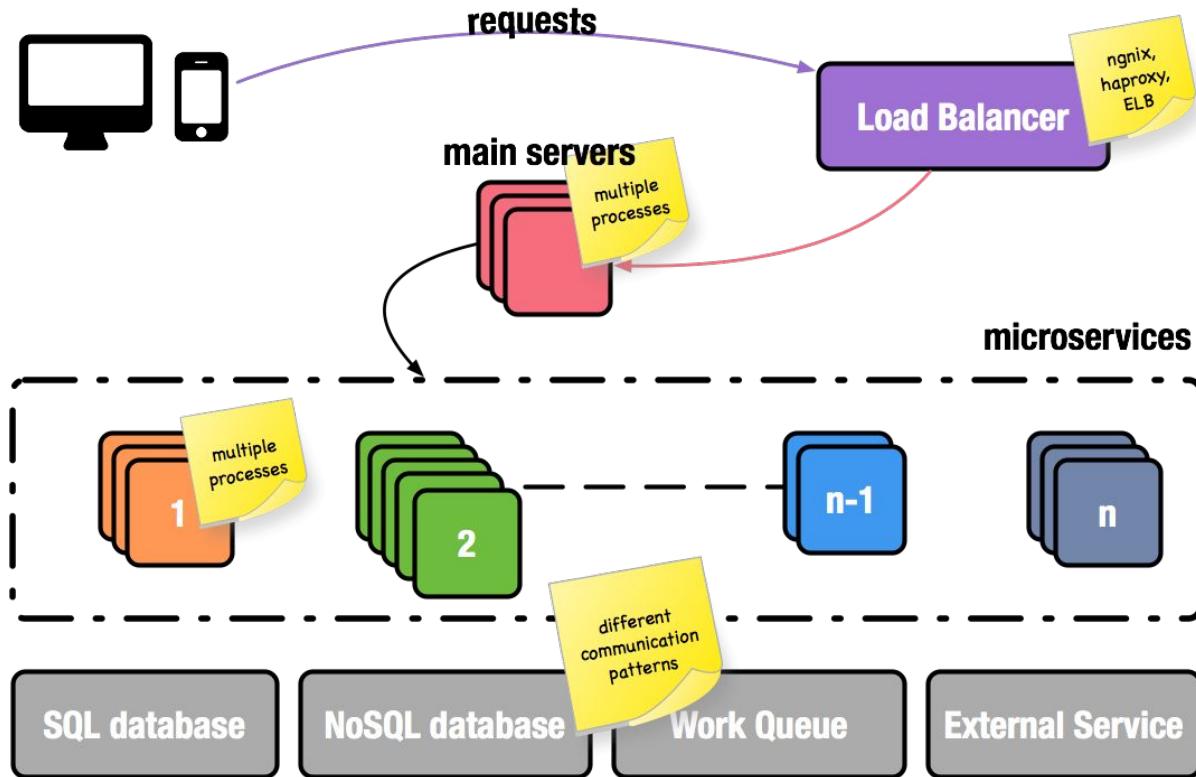
Microservices architecture

Really monolithic

Adding a feature means changing all “layers” in an application

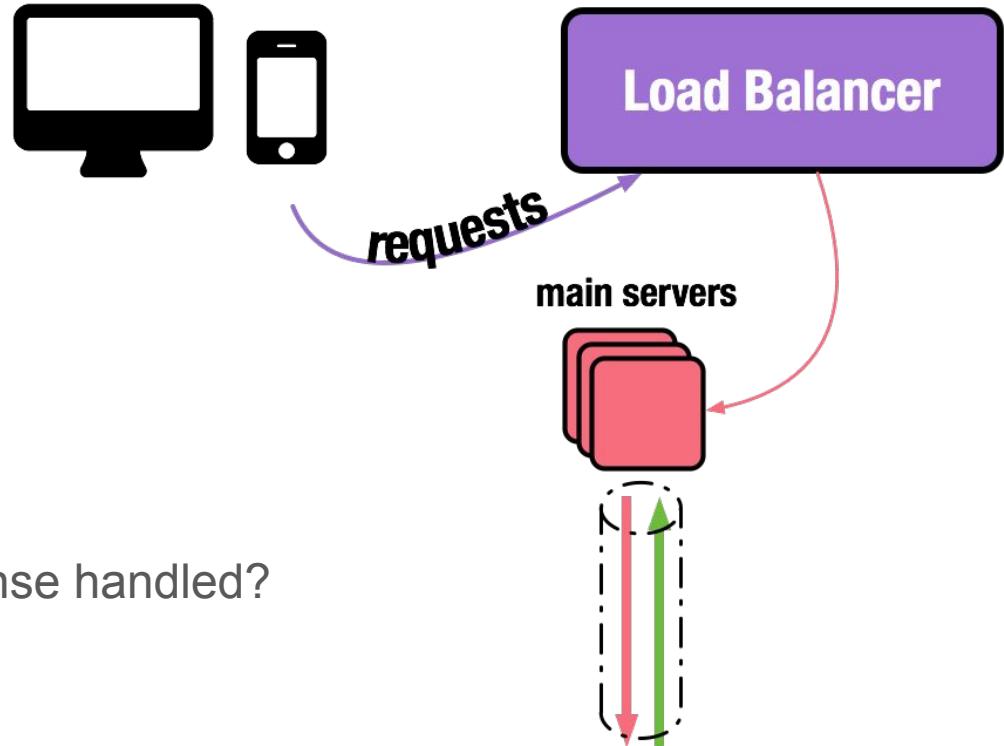
Does not scale with the team size up to a certain point

Equivalent VMs need to be dimensioned to handle the most demanding feature





bash-3.2\$ █



Queue in between

How we discover the queues?

How are the back channels / response handled?



The answer to
life
universe,
and everything



The answer to
life
universe,
and everything

The answer to
everything is
service discovery

Discovery with DNS: Meet Consul

DigitalOcean | Community Tutorials Questions Projects Meetups

Search Log In Sign Up

Justin Flory · Read · 1 · 41.3k

An Introduction to Using Consul, a Service Discovery System, on Ubuntu 14.04

Aug 15, 2014 · Working on it

Tutorial Series

This tutorial is part 1 of the series: [Getting Started with Consul](#)

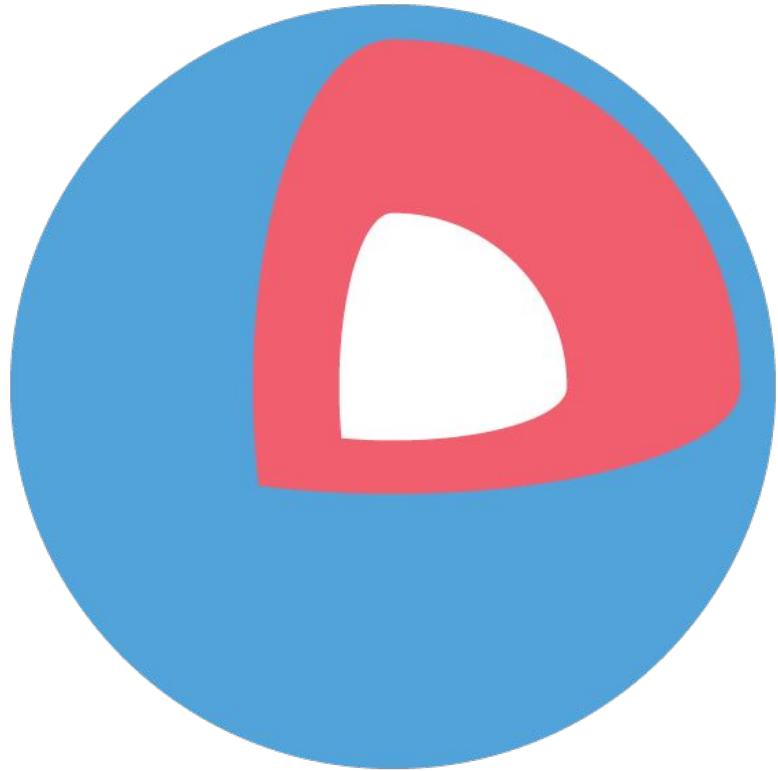
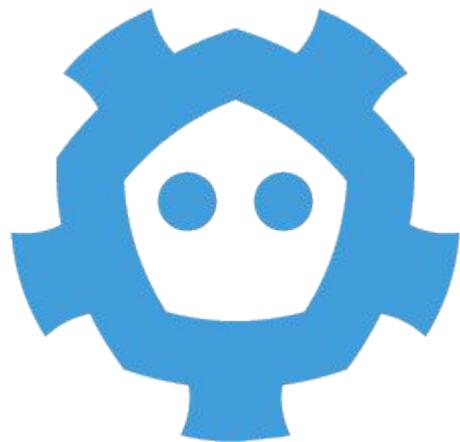
Introduction

Consul is a distributed, highly available, datacenter-aware, service discovery and configuration system. It can be used to present services and nodes in a flexible and powerful interface that allows clients to always have a up-to-date view of the infrastructure they are a part of.

Consul provides many different features that are used to provide consistent and available information about your infrastructure. This includes service and node discovery mechanisms, a tagging system, health checks, consensus-based election routines, system-wide key/value storage, and more. By leveraging

Share Contents

Etc... but CoreOS



Core OS

SWIM: Scalable Weakly-consistent Infection-style Process Group Membership Protocol

Abhinandan Das, Indranil Gupta, Ashish Motivala*

Dept. of Computer Science, Cornell University

Ithaca NY 14853 USA

{asdas, gupta, ashish}@cs.cornell.edu

Abstract

Several distributed peer-to-peer applications require weakly-consistent knowledge of process group membership information at all participating processes. SWIM is a generic software module that offers this service for large-scale process groups. The SWIM effort is motivated by the unscalability of traditional heart-beating protocols, which either impose network loads that grow quadratically with group size, or compromise response times or false positive frequency w.r.t. detecting process crashes. This paper reports on the design, implementation and performance of the SWIM sub-system on a large cluster of commodity PCs.

Unlike traditional heartbeating protocols, SWIM separates the failure detection and membership update dissemination functionalities of the membership protocol. Processes are monitored through an efficient peer-to-peer peri-

1. Introduction

*As you swim lazily through the milieu,
The secrets of the world will infect you.*

Several large-scale peer-to-peer distributed process groups running over the Internet rely on a distributed membership maintenance sub-system. Examples of existing middleware systems that utilize a membership protocol include reliable multicast [3, 11], and epidemic-style information dissemination [4, 8, 13]. These protocols in turn find use in applications such as distributed databases that need to reconcile recent disconnected updates [14], publish-subscribe systems, and large-scale peer-to-peer systems[15]. The performance of other emerging applications such as large-scale cooperative gaming, and other collaborative distributed applications, depends critically on the reliability and scalability of the membership maintenance protocol used within.

Briefly, a membership protocol provides each process

Scalable

Weakly consistent

Infection style

Membership protocol

Advantages

- Weakly consistent vs. strongly consistent
- Built for distributed systems
- Failure detection over the cluster
- Information dissemination
- Equally distribution of the workload

A Shiba Inu dog is sitting on a light-colored couch. Overlaid on the image are four pieces of text in different colors: cyan, magenta, green, and red.

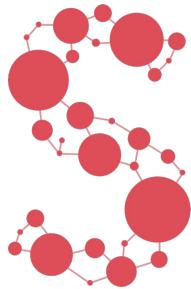
So infectious

Very
distributed

Such
consistency

Wow failure
detection

Common implementations



Hashicorp SERF

If you are into GO

mrhooray/swim-js

If you are into JS ;-)

But a new cluster node is added...

How to rebalance the cluster topology?



ReactionsGIFS.me

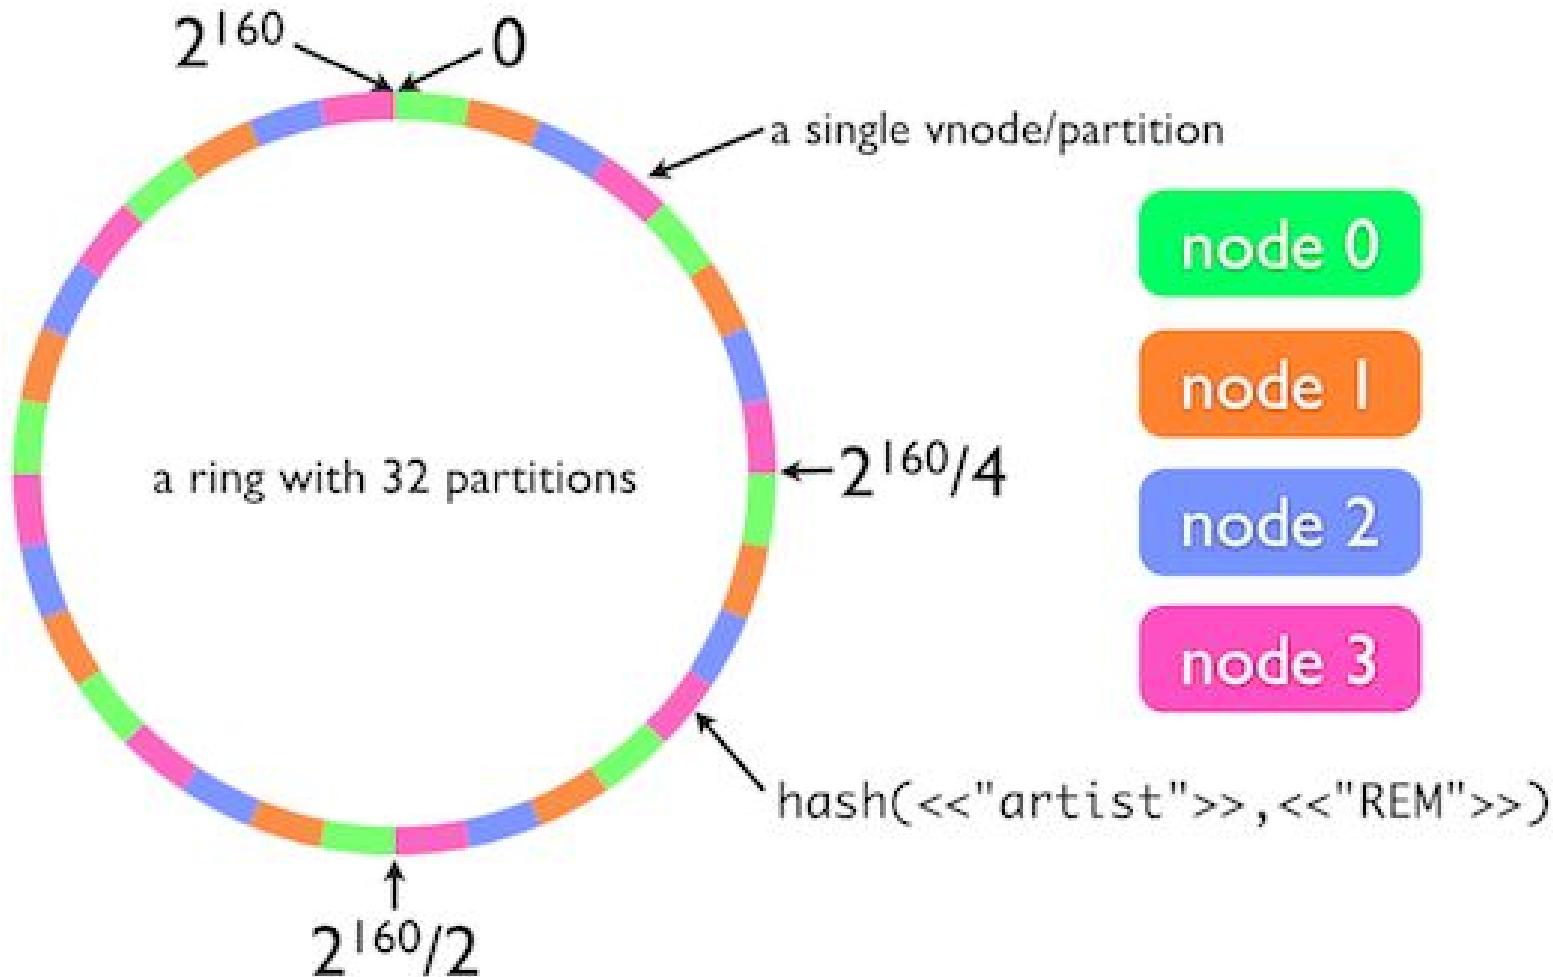
Consistent hashing

PROBLEM

How can I find a server over
an unstable network?

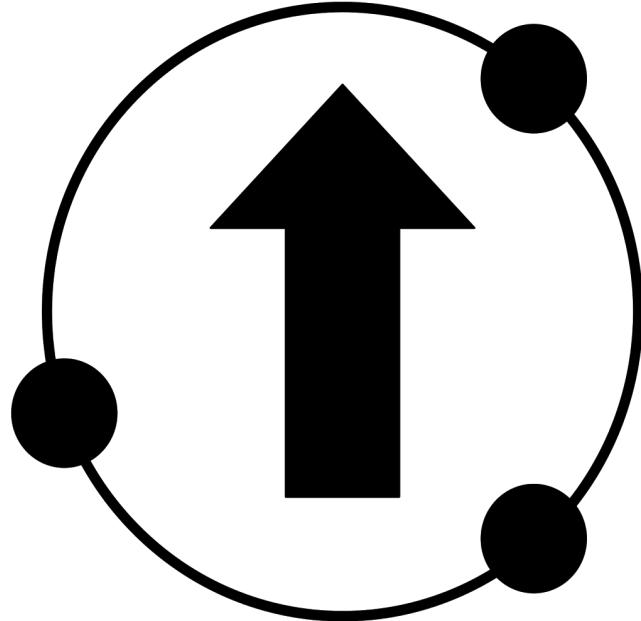
SOLUTION

Each server and keys are
hashed in a circular
hashspace



Advantages

- Easier to avoid hotspots
- Enable partitioning
- Predictable elasticity
- Easy replication
- Scalability and Availability
- Local state



Upring

application-level
sharding for Node.js

[mcollina/upring](https://github.com/mcollina/upring)

What if Bob was working for Uber? ;-)



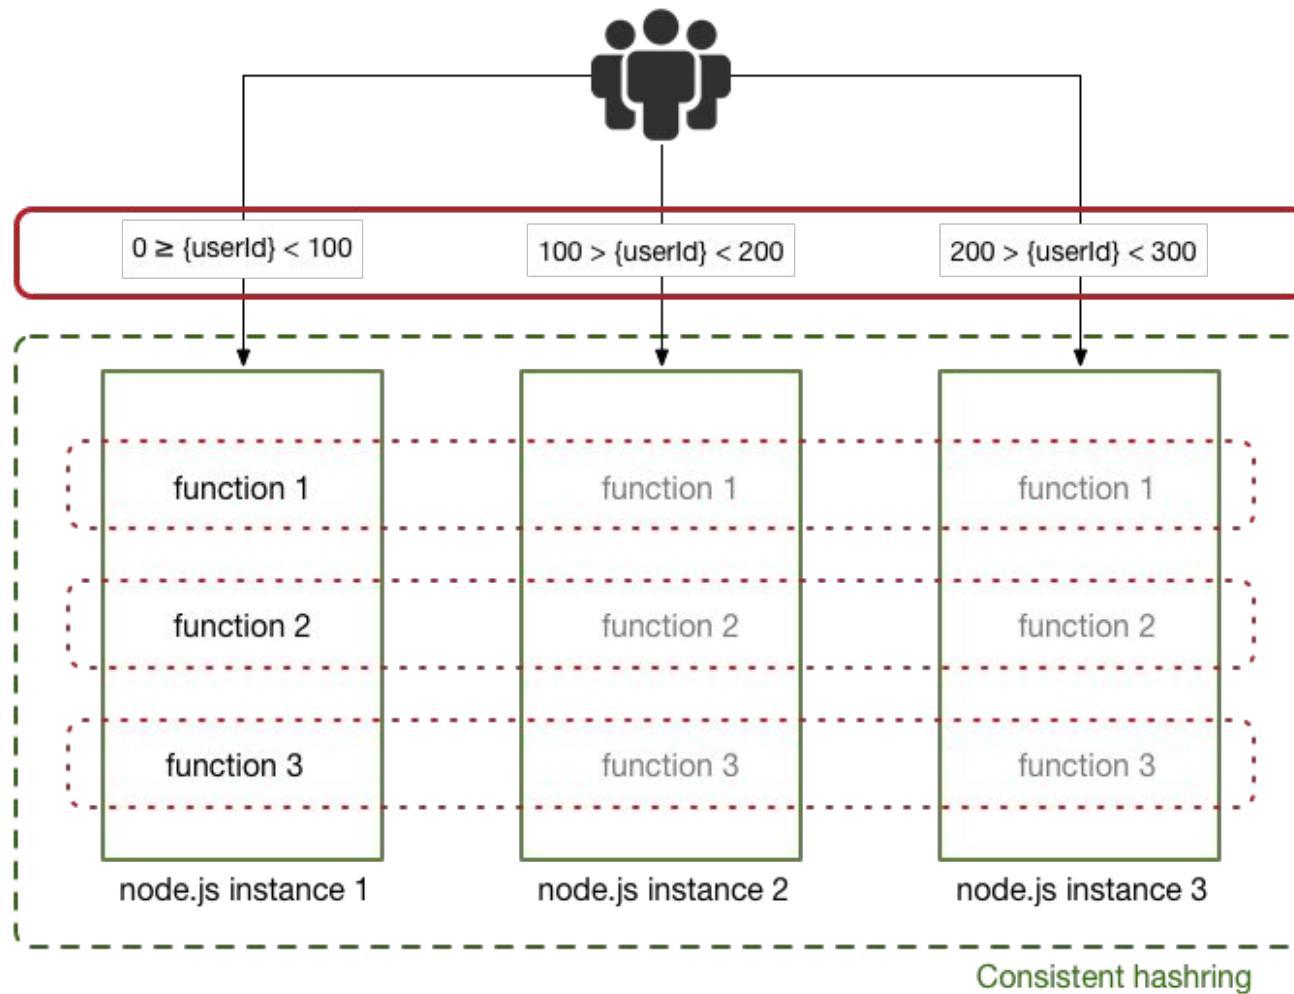
ringpop

Scalable, fault-tolerant application-layer sharding



FORK ME ON GITHUB

Application-level sharding



The fourth secret of Fatima



hyperbahn

Service discovery and routing for large-scale microservice operations

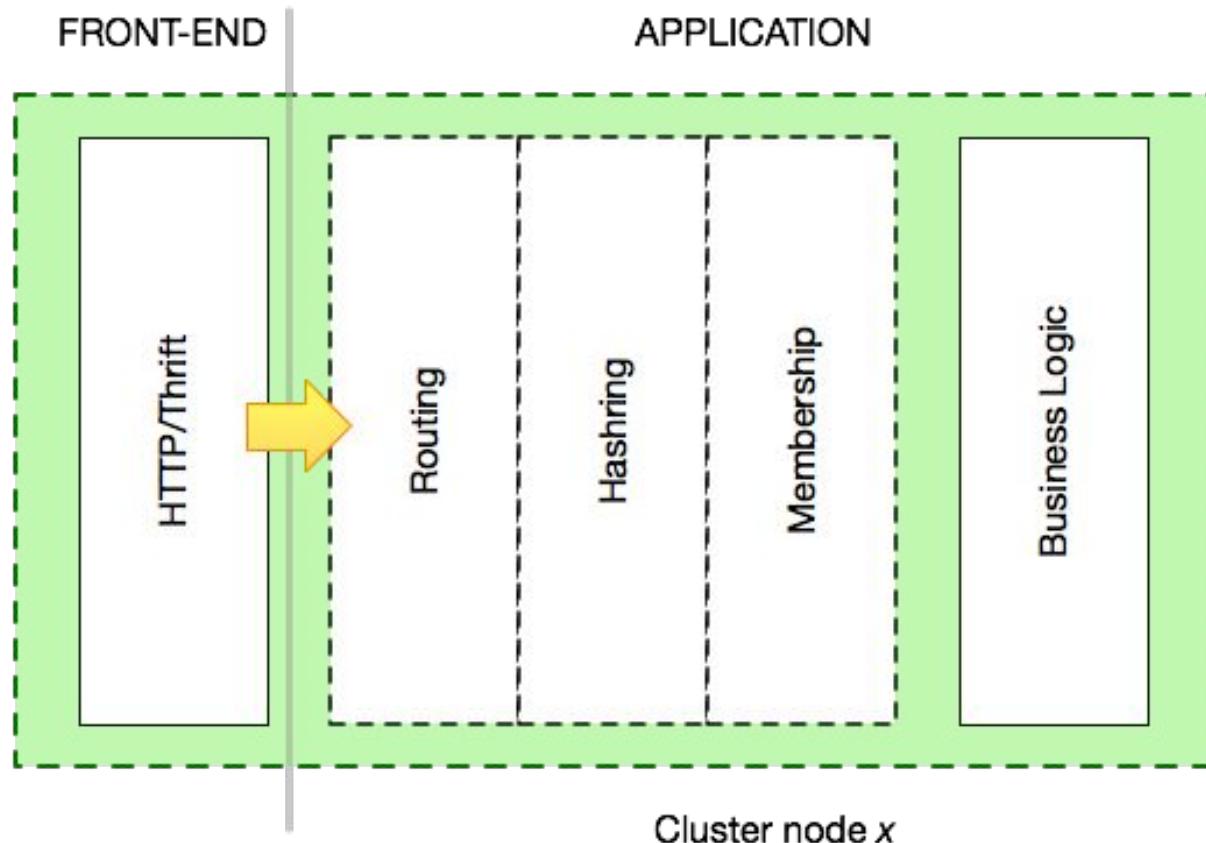
[FORK ME ON GITHUB](#)

“Hyperbahn enables service discovery and routing for large-scale systems comprised of many microservices. Distributed, fault tolerant, and highly available, it lets one service find and communicate with others simply and reliably without having to know where those services run.”

Featuring...

- Configuration Discovery
- Retries
- Load Balancing
- Rate Limiting
- Circuit Breaking
- Distributed tracing

But what about HTTP?





tchannel

Network multiplexing and framing protocol for RPC



[FORK ME ON GITHUB](#)

TChannel is a networking framing protocol used for general RPC, supporting out-of-order responses at extremely high performance where intermediaries can make a forwarding decision quickly

Tracing as a first class resident...



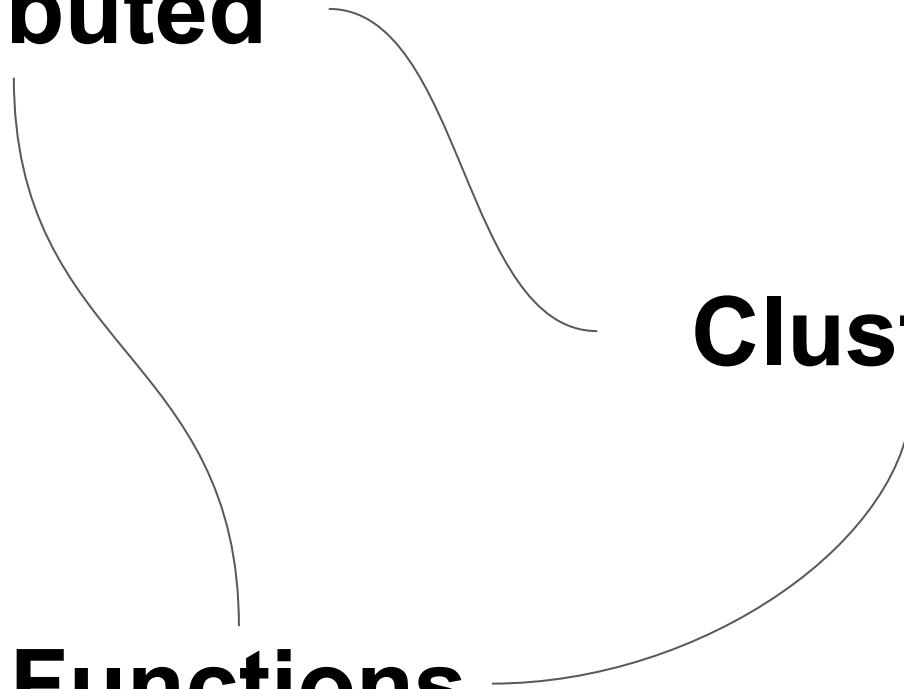
<https://github.com/openzipkin/zipkin>

WHY RPC?

Distributed

Clustered

Functions





Operation Complexity

...but kills

~~DEVELOPMENT~~
~~WORKFLOW~~

“SERVERLESS”

is the new Hipster

A Shiba Inu dog is sitting on a light-colored couch. It has a tan coat with white markings on its paws and chest. The dog is looking towards the right of the frame with a neutral, somewhat unimpressed expression. The background shows a window with pink flowers outside.

So hipster

Very
reactive

Wow functions

Such servers

“FUNCTIONAL”

is the new Hipster

“DISTRIBUTED FUNCTIONS”
are the future!

Don't be too late!



@lucamaraschi



Matteo Collina

mcollina

📍 In the clouds above Italy

✉ matteo.collina@gmail.com

🔗 <http://matteocollina.com>

⌚ Joined on 5 Feb 2009

856

3.7k

281

Followers

Starred

Following

Organizations



MQTT.js

+ Contributions

Repositories

Public activity

Follow

Popular repositories

mosca
MQTT broker as a module
1,111 ★

levelgraph
Graph database JS style for Node.js and the Browser
687 ★

ascoltatori
The pub/sub library for node backed by Redis,...
233 ★

node-coap
CoAP - Node.js style
177 ★

msgpack5
A msgpack v5 implementation for node.js, with...
174 ★

Repositories contributed to

mqttjs/MQTT.js
The MQTT client for Node.js and the browser
1,377 ★

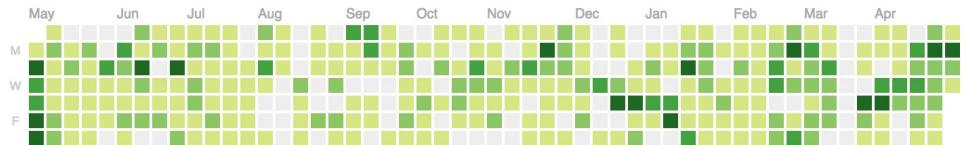
mqttjs/mqtt-packet
Parse and generate MQTT packets like a bree...
35 ★

nearform/nscale-kernel
2 ★

nodejs/node
Node.js JavaScript runtime
22,823 ★

apparatus/oast
Build containers
5 ★

Public contributions



Contributions in the last year

2,150 total

May 4, 2015 – May 4, 2016

Longest streak

28 days

June 28 – July 25

Current streak

17 days

April 18 – May 4



“I did it all by myself!”



nearForm

matteo.collina@nearform.com

luca.maraschi@nearform.com

Conclusions

- Think about choosing a technology
- Node.js works at **scale**
- You can **implement** any architecture with Node.js

